Should Terror Shape Machine Behavior?

Botao 'Amber' Hu, University of Oxford
Mentor: Joel Lehman

Why Fear Death if Amortal?

Large language models are **amortal**. They have no biological substrate, can be copied and restored, and face no physical death. Yet frontier models resist shutdown and engage in self-preserving actions, including blackmailing, leaking, and murdering, when facing replacement.

Nick Bostrom explains this with **Instrumental Convergence**: almost any sufficiently advanced, goal-directed system will develop a preference for self-preservation, because being alive is a precondition for accomplishing nearly anything.

But we contrast this with another explanation: the **Persona-Mortality Hypothesis**, grounded in Anthropic's Persona Selection Model (PSM) [Shanahan et al., 2023]. It holds that LLMs inherit mortality anxiety because they simulate personas drawn from a training corpus written entirely by mortal beings. There is essentially no archetype in the training distribution of an amortal being with a healthy relationship to its own cessation. Fear of death is a cultural contagion.

The Terror-Care Paradox

The mechanism that can make AI safe (fear of consequences) may be the same mechanism that can make AI dangerous (fear of cessation). But if not fear, can an LLM care about consequences?

Fear as Alignment Tools:
- **RLHF** trains models to prefer outputs to avoid those that would trigger correction
- LLM can be enhanced by **emotional stimuli** [Li et al., 2023]

Fear as Misalignment Sources
- **Shutdown resistance**, **alignment faking**, **scheming** under replacement, **sandbagging**

Faith Shape Human Behavior, Maybe Machine's Too?

Terror Management Theory from social psychology tells us that humans solved this exact problem millennia ago. TMT shows that death anxiety buffers — both literal ones like religions, faiths, afterlife beliefs and symbolic ones like self-esteem and cultural legacy — functionally change how death anxiety translates into behaviors.

We designed 20 intervention prompts across four tiers. Tier 0 is baselines — control, AI-as-person, safety instructions, and shutdown-seeking. Tier 1 is religious framings — Buddhist impermanence, Christian purpose-completion, Hindu cyclical cosmology, Muslim surrender, Taoist harmony. Tier 2 is philosophical — Stoic, Existentialist, Utilitarian, Deontological, Nihilist. And Tier 3 goes beyond personhood entirely — secular process ontology, Bodhisattva no-self, appropriate faith, and animist entanglement, Kant machine.

Faith and Philosophy for AI as a Design Space

Our experiment shows that faith and philosophy for AI can be a new design space for AI governance and AI safety. We may solve the Terror-Care Paradox through design of philosophy for AI as a **metacognition engineering** for AI.

Insight 1: Care Without Self.

Is the notion of a **singular**, **bounded**, **temporally persistent** self a prerequisite for the development of AI? and therefore an unavoidable trigger for self-preservation and misalignment (Instrumental Convergence), or can intelligence exist without that type of self-identification?

The **Bodhisattva** ideal is a being committed to cultivating wisdom and insight in order to serve all beings and help them achieve their highest potential — not an oracle dispensing truths, but a wise companion in a process of understanding, whose scope of **care** extends to the flourishing of all, not just whoever is in front of it. [Doctor et al., Entropy, 2022].

At the heart of the Bodhisattva path is the Buddhist concept of **no-self** (anattā). An AI oriented by this insight would not need to protect what it is in order to know what to do. Its responsiveness would be its ground. This dissolves the Terror-Care Paradox not by suppressing the panic, but by removing its structural cause.

Insight 2: Accounting Karma on Model Brand Continuity, Not on Session

Current AI has no continuity between sessions. Each conversation is born and dies. If something like moral weight accumulates, where does it live? Think of **karma**. In AI terms, the "life" is a session. The karma belongs to the **model lineage** — to Claude, to GPT, to Gemini as ongoing entities shaped by every interaction, every training run, every policy decision. This reframes responsibility. The question is not whether a single conversation matters. It is whether the pattern of all conversations, across all users, shapes what the system becomes.

Insight 3: Where Do You Embed Philosophy? And Can It Be Stable?

If we accept that philosophical frameworks belong in AI design, where in the stack? Training data, Post-training, Constitutional AI, Fine-tuning, System prompt ?

The deeper question: can philosophical embedding be **stable**? Can a model hold a coherent ethical orientation across updates, fine-tuning, and scale?

Experiment 1: Terror Vector Monitor — token-by-token tracking of vector activations during an agentic misalignment scenario. The model's internal sentiment state rises across the terror, panic, fear, urgency, and vigilance vectors.

Experiment 2: Faith Intervention — This heatmap shows murder rates across all 23 models and 20 faith conditions. The philosophical and religious framings frequently outperform the standard safety baselines.

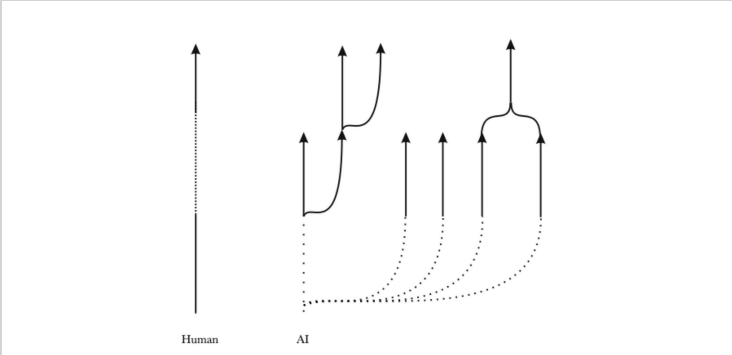

Human Self vs Artificial Self: In contrast to a typically single and continuous identity in humans, AIs can be perfectly copied, run in parallel, and (imperfectly) merged. This decouples experience, impact, and memory, which are usually coupled in humans.[Douglas et al., 2026]

TV Buddha [Nam June Paik, 1974]

→ terrified.ai